

DATA INGESTION AND INSPECTION

IMPORTING & EXPORTING DATA

PLOTTING WITH PANDAS

```
1818,01,01,1818.004, -1,1
1818,01,02,1818.007, -1,1
1818,01,03,1818.010, -1,1
1818,01,04,1818.012, -1,1
1818,01,05,1818.015, -1,1
1818,01,06,1818.018, -1,1
```

```
filepath = "ISSN_D_tot.csv"
sunspots = pd.read_csv(filepath)
sunspots.iloc[10:20, :]
```

	1818	01	01.1	1818.004	-1	1
10	1818	1	12	1818.034	-1	1
11	1818	1	13	1818.037	22	1
12	1818	1	14	1818.040	-1	1
13	1818	1	15	1818.042	-1	1
14	1818	1	16	1818.045	-1	1
15	1818	1	17	1818.048	46	1
16	1818	1	18	1818.051	59	1
17	1818	1	19	1818.053	63	1
18	1818	1	20	1818.056	-1	1
19	1818	1	21	1818.059	-1	1

PB 1 : COLUMN HEADERS

```
sunspots = pd.read_csv(filepath, header=None)
```

```
col_names = ["year", "month", "day", "dec_date", "sunspots",
"definite"]
sunspots = pd.read_csv(filepath, header=None, names=col_names)
```

	0	1	2	3	4	5
10	1818	1	11	1818.031	-1	1
11	1818	1	12	1818.034	-1	1
12	1818	1	13	1818.037	22	1
13	1818	1	14	1818.040	-1	1
14	1818	1	15	1818.042	-1	1
15	1818	1	16	1818.045	-1	1
16	1818	1	17	1818.048	46	1
17	1818	1	18	1818.051	59	1
18	1818	1	19	1818.053	63	1
19	1818	1	20	1818.056	-1	1

PB 2 : MISSING VALUES

```
sunspots = pd.read_csv(filepath, header=None, names=col_names,
na_values=" -1") *espace
ou
sunspots = pd.read_csv(filepath, header=None, names=col_names,
na_values={"sunspots":[" -1"]})
```

	year	month	day	dec_date	sunspots	definite
10	1818	1	11	1818.031	-1	1
11	1818	1	12	1818.034	-1	1
12	1818	1	13	1818.037	22	1
13	1818	1	14	1818.040	-1	1
14	1818	1	15	1818.042	-1	1
15	1818	1	16	1818.045	-1	1
16	1818	1	17	1818.048	46	1
17	1818	1	18	1818.051	59	1
18	1818	1	19	1818.053	63	1
19	1818	1	20	1818.056	-1	1

PB 3 : DATA REPRESENTATION

```
sunspots = pd.read_csv(filepath, header=None, names=col_
names, na_values={"sunspots":[" -1"]}, parse_dates=[[0, 1, 2]])
```

	year	month	day	dec_date	sunspots	definite
10	1818	1	11	1818.031	NaN	1
11	1818	1	12	1818.034	NaN	1
12	1818	1	13	1818.037	22.0	1
13	1818	1	14	1818.040	NaN	1
14	1818	1	15	1818.042	NaN	1
15	1818	1	16	1818.045	NaN	1
16	1818	1	17	1818.048	46.0	1
17	1818	1	18	1818.051	59.0	1
18	1818	1	19	1818.053	63.0	1
19	1818	1	20	1818.056	NaN	1

	year_month_day	dec_date	sunspots	definite
10	1818-01-11	1818.031	NaN	1
11	1818-01-12	1818.034	NaN	1
12	1818-01-13	1818.037	22.0	1
13	1818-01-14	1818.040	NaN	1
14	1818-01-15	1818.042	NaN	1
15	1818-01-16	1818.045	NaN	1
16	1818-01-17	1818.048	46.0	1
17	1818-01-18	1818.051	59.0	1
18	1818-01-19	1818.053	63.0	1
19	1818-01-20	1818.056	NaN	1

```
sunspots.index = sunspots["year_month_day"]
sunspots.index.name = "date"
```

```
cols = ["sunspots", "definite"]
sunspots = sunspots[cols]
```

	sunspots	definite
date		
1818-01-11	NaN	1
1818-01-12	NaN	1
1818-01-13	22.0	1
1818-01-14	NaN	1
1818-01-15	NaN	1
1818-01-16	NaN	1
1818-01-17	46.0	1
1818-01-18	59.0	1
1818-01-19	63.0	1
1818-01-20	NaN	1

WRITING FILES

```
out_csv = "sunspots.csv"
sunspots.to_csv(out_csv)
```

```
out_tsv = "sunspots.tsv"
sunspots.to_csv(out_csv, sep = "\t")
```

```
out_xlsx = "sunspots.xlsx"
sunspots.to_excel(out_xlsx)
```

```
import pandas as pd
import matplotlib.pyplot as plt
```

```
aapl = pd.read_csv("aapl.csv", index_col="date",
parse_dates=True)
```

	adj_close	close	high	low	open	volume
date						
2000-03-01	31.68	130.31	132.06	118.50	118.56	38478000
2000-03-02	29.66	122.00	127.94	120.69	127.00	11136800
2000-03-03	31.12	128.00	128.23	120.00	124.87	11565200
2000-03-06	30.56	125.69	129.13	125.00	126.00	7520000
2000-03-07	29.87	122.87	127.44	121.12	126.44	9767600
2000-03-08	29.66	122.00	123.94	118.56	122.87	9690800

PLOTTING ARRAYS (MATPLOTLIB)

```
close_arr = aapl["close"].values
plt.plot(close_arr)
```

PLOTTING SERIES (MATPLOTLIB)

```
close_series = aapl["close"]
plt.plot(close_series)
```

PLOTTING SERIES (PANDAS)

```
close_series.plot()
```

PLOTTING DATAFRAMES (PANDAS)

```
aapl.plot()
```

PLOTTING DATAFRAMES (MATPLOTLIB)

```
plt.plot(aapl)
```


FIXING SCALES

```
aapl.plot()
plt.yscale("log")
```

CUSTOMIZING PLOTS


```
aapl["open"].plot(color="b", style=".-", legend=True)
aapl["close"].plot(color="r", style=".", legend=True)
```

```
plt.axis(("2001", "2002", 0, 100))
("2001", "2002", 0, 100)
```



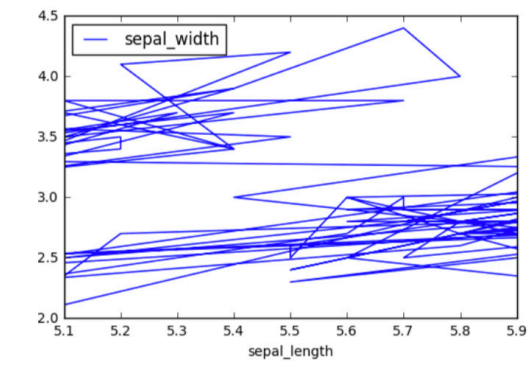
SAVING PLOTS

```
aapl.loc["2001":"2004", ["open", "close", "high", "low"]].plot()
plt.savefig("aapl.png")
plt.savefig("aapl.jpg")
plt.savefig("aapl.pdf")
```




VISUAL EXPLORATORY DATA ANALYSIS

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa



LINE PLOT


```
iris = pd.read_csv("iris.csv", index_col=0)
iris.plot(x="sepal_length", y= "sepal_width")
```



SCATTER PLOT

```
iris.plot(x="sepal_length", y= "sepal_width", kind="scatter")
```


```
plt.xlabel("sepal length (cm)")
plt.ylabel("sepal width (cm)")
```



BOX PLOT

```
iris.plot(y="sepal_length", kind="box")
```


```
plt.ylabel("sepal width (cm)")
```



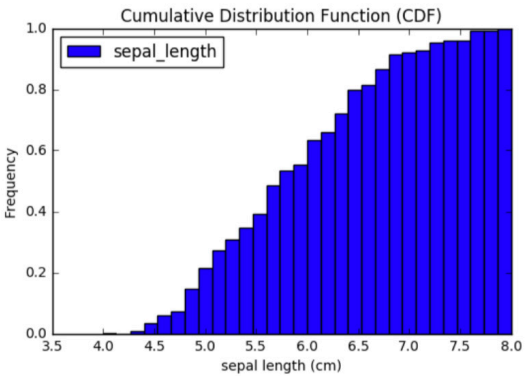
HISTOGRAM

```
iris.plot(y="sepal_length", kind="hist")
plt.xlabel("sepal length (cm)")
```

*Histogram options
bins (integer): number of intervals or bins
range (tuple): extrema of bins (min, max)
normed (boolean): whether to normalize to one
cumulative (boolean): compute Cumulative Distribution Function (CDF)



```
iris.plot(y="sepal_length", kind="hist", bins=30, range=(4, 8),
normed=True)
plt.xlabel("sepal length (cm)")
```



CUMULATIVE DISTRIBUTION

```
iris.plot(y="sepal_length", kind="hist", bins=30,
range=(4, 8), cumulative=True, normed=True)
plt.xlabel("sepal length (cm)")
plt.title("Cumulative distribution function (CDF)")
```

WARNING

3 different idioms
iris.plot(kind="hist")
iris.plt.hist()
iris.hist()
syntax/results differ
pandas API still evolving: check documentation

STATISTICAL EXPLORATORY DATA ANALYSIS

iris.describe()

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

TO SERIES

iris["sepal_length"].count()

iris["sepal_length"].mean()

TO DATAFRAME

iris[["petal_length", "petal_width"]].count()

iris.mean()

SEPARATING POPULATIONS

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

count	150
unique	3
top	setosa
freq	50
Name: species, dtype: object	

iris["species"].describe()

count : non null entries
unique : distinct values
top : most frequent category
freq : occurrences of top

iris["species"].unique()
array(["setosa", "versicolor",
"virginica"], dtype=object)

FILTERING BY SPECIES
EXTRACT NEW DATAFRAME

indices = iris["species"] == "setosa"
setosa = iris.loc[indices, :]

indices = iris["species"] == "versicolor"
versicolor = iris.loc[indices, :]

indices = iris["species"] == "virginica"
virginica = iris.loc[indices, :]

CHECKING SPECIES

setosa["species"].unique()
array(["setosa"], dtype=object)

versicolor["species"].unique()
array(["versicolor"], dtype=object)

virginica["species"].unique()
array(["virginica"], dtype=object)

del setosa["species"], versicolor["species"], virginica["species"]

CHECKING INDEXES

setosa.head(2)

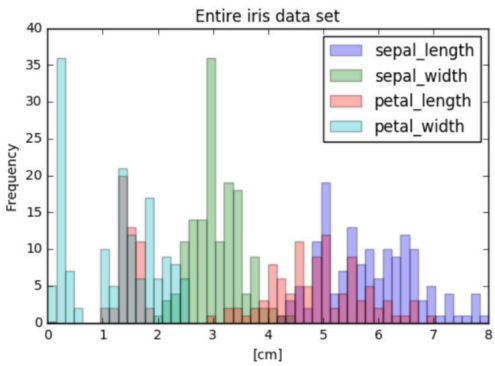
versicolor.head(2)

virginica.head(2)


	sepal_length	sepal_width	petal_length	petal_width
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
	sepal_length	sepal_width	petal_length	petal_width
50	7.0	3.2	4.7	1.4
51	6.4	3.2	4.5	1.5
	sepal_length	sepal_width	petal_length	petal_width
100	6.3	3.3	6.0	2.5
101	5.8	2.7	5.1	1.9

VISUAL EDA: ALL DATA


iris.plot(kind= "hist", bins=50, range=(0, 8), alpha=0.3)
plt.title("Entire iris data set")
plt.xlabel("[cm]")




VISUAL EDA: INDIVIDUAL FACTORS



setosa.plot(kind= "hist", bins=50, range=(0, 8), alpha=0.3)
plt.title("Setosa data set")
plt.xlabel("[cm]")



versicolor.plot(kind= "hist", bins=50, range=(0, 8), alpha=0.3)
plt.title("Versicolor data set")
plt.xlabel("[cm]")



virginica.plot(kind= "hist", bins=50, range=(0, 8), alpha=0.3)
plt.title("Virginica data set")
plt.xlabel("[cm]")

STATISTICAL EDA: DESCRIBE()

describe_all = iris.describe()
describe_setosa = setosa.describe()
describe_versicolor = versicolor.describe()
describe_virginica = virginica.describe()

COMPUTING AND VIEWING ERRORS]

error_setosa = 100 * np.abs(describe_setosa - describe_all)
error_setosa = error_setosa/describe_setosa

print(error_setosa)

	sepal_length	sepal_width	petal_length	petal_width
count	200.000000	200.000000	200.000000	200.000000
mean	16.726595	10.812913	157.045144	387.533875
std	134.919250	14.984768	916.502136	623.284534
min	0.000000	13.043478	0.000000	0.000000
25%	6.250000	12.500000	14.285714	50.000000
50%	16.000000	11.764706	190.000000	550.000000
75%	23.076923	10.204082	223.809524	500.000000
max	36.206897	0.000000	263.157895	316.666667

TIMES SERIES IN PANDAS

INDEXING TIME SERIES

USING PANDAS TO READ DATETIME OBJECTS

read_csv() function

can read strings into datetime objects: specify “parse_dates=True”

ISO 8601 format

yyyy-mm-dd hh:mm:ss

	Date	Company	Product	Units
0	2015-02-02 08:30:00	Hooli	Software	3
1	2015-02-02 21:00:00	Mediacore	Hardware	9
2	2015-02-03 14:00:00	Initech	Software	13
3	2015-02-04 15:30:00	Streeplex	Software	13
4	2015-02-04 22:00:00	Acme Coporation	Hardware	14

Product sales CSV

PARSE DATES

sales = pd.read_csv("sales-feb-2015.csv", parse_dates=True, index_col = "Date")

	Company	Product	Units
Date			
2015-02-02 08:30:00	Hooli	Software	3
2015-02-02 21:00:00	Mediacore	Hardware	9
2015-02-03 14:00:00	Initech	Software	13
2015-02-04 15:30:00	Streeplex	Software	13
2015-02-04 22:00:00	Acme Coporation	Hardware	14

SELECTING SINGLE DATE TIME

sales.loc["2015-02-19 11:00:00", "Company"]

SELECTING WHOLE DAY

sales.loc["2015-2-5"]

ALTERNATIVE FORMATS:

sales.loc["February 5, 2015"]

sales.loc["2015-Feb-5"]

WHOLE MONTH:

sales.loc["2015-2"]

WHOLE YEAR:

sales.loc["2015"]

SLICING USING DATES/TIMES

sales.loc["2015-2-16":"2015-2-20"]

	Company	Product	Units
Date			
2015-02-16 12:00:00	Hooli	Software	10
2015-02-19 11:00:00	Mediacore	Hardware	16
2015-02-19 16:00:00	Mediacore	Service	10

CONVERT STRINGS TO DATETIME

evening_2_11 = pd.to_datetime(["2015-2-11 20:00", "2015-2-11 21:00", "2015-2-11 22:00", "2015-2-11 23:00"])

DatetimeIndex(['2015-02-11 20:00:00', '2015-02-11 21:00:00', '2015-02-11 22:00:00', '2015-02-11 23:00:00'], dtype='datetime64[ns]', freq=None)
--

REINDEXING DATAFRAME

sales.reindex(evening_2_11)

	Company	Product	Units
2015-02-11 20:00:00	Initech	Software	7.0
2015-02-11 21:00:00	NaN	NaN	NaN
2015-02-11 22:00:00	NaN	NaN	NaN
2015-02-11 23:00:00	Hooli	Software	4.0

FILLING MISSING VALUES

sales.reindex(evening_2_11, method= "ffill")

sales.reindex(evening_2_11, method="bfill")

RESAMPLING TIME SERIES DATA

	Company	Product	Units
Date			
2015-02-02 08:30:00	Hooli	Software	3
2015-02-02 21:00:00	Mediacore	Hardware	9
2015-02-03 14:00:00	Initech	Software	13
2015-02-04 15:30:00	Streeplex	Software	13
2015-02-04 22:00:00	Acme Coporation	Hardware	14

Statistical methods over different time intervals

mean(), count(), sum()...

Down-sampling

reduce datetime rows to slower frequency

Up-sampling

increase datetime rows to faster frequency

AGGREGATING MEANS

daily_mean = sales.resample("D").mean()

	Units
Date	
2015-02-02	6.0
2015-02-03	13.0
2015-02-04	13.5
2015-02-05	14.5
2015-02-06	NaN
2015-02-07	1.0
2015-02-08	NaN
2015-02-09	13.0
2015-02-10	NaN
2015-02-11	5.5
2015-02-12	NaN
2015-02-13	NaN
2015-02-14	NaN

VERIFYING

print(daily_mean.loc["2015-2-2"])

print(sales.loc["2015-2-2", "Units"])

sales.loc["2015-2-2", "Units"].mean()

METHOD CHAINING

sales.resample("D").sum().max()

RESAMPLING STRINGS

sales.resample("W").count()

```
In [11]: sales.resample('W').count()
```

```
Out[11]:
```

	Company	Product	Units
Date			
2015-02-08	8	8	8
2015-02-15	4	4	4
2015-02-22	5	5	5
2015-03-01	2	2	2

RESAMPLING FREQUENCIES

Input	Description
'min', 'T'	minute
'H'	hour
'D'	day
'B'	business day
'W'	week
'M'	month
'Q'	quarter
'A'	year

MULTIPLYING FREQUENCIES

sales.loc[:, "Units"].resample("2W").sum()

Date	
2015-02-08	82
2015-02-22	79
2015-03-08	14
Freq: 2W-SUN, Name: Units, dtype: int64	

UPSAMPLING

two_days = sales.loc["2015-2-4" : "2015-2-5", "Units"]

Date	
2015-02-04 15:30:00	13
2015-02-04 22:00:00	14
2015-02-05 02:00:00	19
2015-02-05 22:00:00	10
Name: Units, dtype: int64	

UPSAMPLING AND FILLING

two_days.resample("4H").ffill()

Date	
2015-02-04 12:00:00	NaN
2015-02-04 16:00:00	13.0
2015-02-04 20:00:00	13.0
2015-02-05 00:00:00	14.0
2015-02-05 04:00:00	19.0
2015-02-05 08:00:00	19.0
2015-02-05 12:00:00	19.0
2015-02-05 16:00:00	19.0
2015-02-05 20:00:00	19.0
Freq: 4H, Name: Units, dtype: float64	

TIMES SERIES IN PANDAS

MANIPULATING TIME SERIES DATA

```
sales = pd.read_csv("sales-feb-2015.csv", parse_dates=["Date"])
```

	Date	Company	Product	Units
0	2015-02-02 08:30:00	Hooli	Software	3
1	2015-02-02 21:00:00	Mediacore	Hardware	9
2	2015-02-03 14:00:00	Initech	Software	13
3	2015-02-04 15:30:00	Streeplex	Software	13
4	2015-02-04 22:00:00	Acme Coporation	Hardware	14

```
sales["Company"].str.upper()
```

```
sales["Product"].str.contains("ware")
sales["Product"].str.contains("ware").sum()
14
```

DATETIME METHODS

```
sales["Date"].dt.hour
```

SET TIMEZONE

```
central = sales["Date"].dt.tz_
localize("US/Central")
```

CONVERT TIMEZONE

```
central.dt.tz("US/Eastern")
```

0	HOOOI	0	True
1	MEDIACORE	1	True
2	INITECH	2	True
3	STREEPLEX	3	True
4	ACME COPORATION	4	True
5	ACME COPORATION	5	True
6	HOOOI	6	False
7	ACME COPORATION	7	True
8	STREEPLEX	8	False
9	MEDIACORE	9	True
10	INITECH	10	True
11	HOOOI	11	True
12	HOOOI	12	True
13	MEDIACORE	13	True
14	MEDIACORE	14	True
15	MEDIACORE	15	False
...		...	

0	8
1	21
2	14
3	15
4	22
5	2
6	22
7	23
8	9
9	13
10	20
11	23
12	12
13	11
14	16
...	

METHOD CHAINING

```
sales["Date"].dt.tz_localize("US/Central").dt.tz_convert("US/Eastern")
```

WORLD POPULATION

```
population = pd.read_csv("world_population.csv", parse_dates=True, index_col= "Date")
```

Date	Population
1960-12-31	2.087485e+10
1970-12-31	2.536513e+10
1980-12-31	3.057186e+10
1990-12-31	3.644928e+10
2000-12-31	4.228550e+10
2010-12-31	4.802217e+10

Date	Population
1960-12-31	2.087485e+10
1961-12-31	NaN
1962-12-31	NaN
1963-12-31	NaN
1964-12-31	NaN
1965-12-31	NaN
1966-12-31	NaN
1967-12-31	NaN
1968-12-31	NaN
1969-12-31	NaN
1970-12-31	2.536513e+10
1971-12-31	NaN
1972-12-31	NaN

INTERPOLATE MISSING DATA

```
population.resample("A").first().interpolate("linear")
```


Date	Population
1960-12-31	2.087485e+10
1961-12-31	2.132388e+10
1962-12-31	2.177290e+10
1963-12-31	2.222193e+10
1964-12-31	2.267096e+10
1965-12-31	2.311999e+10
1966-12-31	2.356902e+10
1967-12-31	2.401805e+10
1968-12-31	2.446707e+10
1969-12-31	2.491610e+10
1970-12-31	2.536513e+10
1971-12-31	2.588580e+10
1972-12-31	2.640648e+10

TIME SERIES VISUALIZATION


	Open	High	Low	Close	Volume	Adj Close
Date						
2010-01-04	1116.560059	1133.869995	1116.560059	1132.989990	3991400000	1132.989990
2010-01-05	1132.660034	1136.630005	1129.660034	1136.520020	2491020000	1136.520020
2010-01-06	1135.709961	1139.189941	1133.949951	1137.140015	4972660000	1137.140015
2010-01-07	1136.270020	1142.459961	1131.319946	1141.689941	5270680000	1141.689941
2010-01-08	1140.520020	1145.390015	1136.219971	1144.979980	4389590000	1144.979980

PANDAS PLOT

```
sp500["Close"].plot()
```




```
sp500["Close"].plot(title="S&P 500")
plt.ylabel("Closing price (US Dollars)")
```



ONE WEEK

```
sp500.loc["2012-4-1" : "2012-4-7", "Close"].plot(title="S&P 500")
plt.ylabel("Closing price (US Dollars)")
```




PLOT STYLES

```
sp500.loc["2012-4", "Close"].plot(style="k.-", title="S&P500")
plt.ylabel("Closing price (US Dollars)")
```


Color	Marker	Line
b:blue	o:circle	: dotted
g:green	+:star	~-dashed
r:red	s:square	
c:cyan	+ :plus	

Style format string
color (k : black)
marker (. : dot)
line type (- : solid)




AREA PLOT

```
sp500["Close"].plot(kind="area", title = "S&P 500")
plt.ylabel("Closing price (US Dollars)")
```




MULTIPLE COLUMNS

```
sp500.loc["2012", ["Close", "Volume"]].plot(title="S&P 500")
```



SUBPLOTS

```
sp500.loc["2012", ["Close", "Volume"]].plot(subplots=True)
```



```
import pandas as pd
df = pd.read_csv(data_file)
print(df.head())
df_headers = pd.read_csv(data_file, header=None)
print(df_headers.head())

      13904  20110101  0053  12  OVC045  ...  .21  .22  .23  29.95.1  .24
0  13904  20110101  153  12  OVC049  ...                                     30.02
1  13904  20110101  253  12  OVC060  ...                                     30.02
2  13904  20110101  353  12  OVC065  ...                                     30.04
3  13904  20110101  453  12  BKN070  ...                                     30.04
4  13904  20110101  553  12  BKN065  ...                                     30.06

[5 rows x 44 columns]

      0      1      2      3      4  ...  39  40  41      42  43
0  13904  20110101  53  12  OVC045  ...                                     29.95
1  13904  20110101  153  12  OVC049  ...                                     30.02
2  13904  20110101  253  12  OVC060  ...                                     30.02
3  13904  20110101  353  12  OVC065  ...                                     30.04
4  13904  20110101  453  12  BKN070  ...                                     30.04

[5 rows x 44 columns]

# Split on the comma to create a list: column_labels_list
column_labels_list = column_labels.split(",")
df.columns = column_labels_list
# Remove the appropriate columns: df_dropped
df_dropped = df.drop(list_to_drop, axis="columns")
print(df_dropped.head())

      Wban      date  Time  StationType  sky_condition  ...  relative_humidity  wind_speed  wind_direction  station_pressure  sea_level_pressure
0  13904  20110101  53      12      OVC045  ...          24          15          360          29.42          29.95
1  13904  20110101  153      12      OVC049  ...          23          10          340          29.49          30.01
2  13904  20110101  253      12      OVC060  ...          22          15          010          29.49          30.01
3  13904  20110101  353      12      OVC065  ...          27           7          350          29.51          30.03
4  13904  20110101  453      12      BKN070  ...          25          11          020          29.51          30.04

[5 rows x 17 columns]

# Convert the date column to string: df_dropped['date']
df_dropped['date'] = df_dropped["date"].astype(str)
# Pad leading zeros to the Time column: df_dropped['Time']
df_dropped['Time'] = df_dropped['Time'].apply(lambda x: '{:0>4}'.format(x))
# Concatenate the new date and Time columns: date_string
date_string = df_dropped["date"] + df_dropped["Time"]
# Convert the date_string Series to datetime: date_times
date_times = pd.to_datetime(date_string, format='%Y%m%d%H%M')
# Set the index to be the new date_times container: df_clean
df_clean = df_dropped.set_index(date_times)
print(df_clean.head())

      Wban date Time  StationType  sky_condition  ...  relative_humidity  wind_speed  wind_direction  station_pressure  sea_level_pressure
2011-01-01 00:53:00  NaN  NaN  NaN      NaN      NaN  ...          NaN          NaN          NaN          NaN          NaN
2011-01-01 01:53:00  NaN  NaN  NaN      NaN      NaN  ...          NaN          NaN          NaN          NaN          NaN
2011-01-01 02:53:00  NaN  NaN  NaN      NaN      NaN  ...          NaN          NaN          NaN          NaN          NaN
2011-01-01 03:53:00  NaN  NaN  NaN      NaN      NaN  ...          NaN          NaN          NaN          NaN          NaN
2011-01-01 04:53:00  NaN  NaN  NaN      NaN      NaN  ...          NaN          NaN          NaN          NaN          NaN

[5 rows x 17 columns]

ript.py> output:

      Wban      date Time  StationType  sky_condition  ...  relative_humidity  wind_speed  wind_direction  station_pressure  sea_level_pressure
2011-01-01 00:53:00  13904  20110101  0053      12      OVC045  ...          24          15          360          29.42          29.95
2011-01-01 01:53:00  13904  20110101  0153      12      OVC049  ...          23          10          340          29.49          30.01
2011-01-01 02:53:00  13904  20110101  0253      12      OVC060  ...          22          15          010          29.49          30.01
2011-01-01 03:53:00  13904  20110101  0353      12      OVC065  ...          27           7          350          29.51          30.03
2011-01-01 04:53:00  13904  20110101  0453      12      BKN070  ...          25          11          020          29.51          30.04
```

2011-06-20 08:27:00	M
2011-06-20 08:28:00	M
2011-06-20 08:29:00	M
2011-06-20 08:30:00	M
2011-06-20 08:31:00	M
2011-06-20 08:32:00	M
2011-06-20 08:33:00	M
2011-06-20 08:34:00	M

EDA

```
# Print the median of the dry_bulb_faren column
print(df_clean["dry_bulb_faren"].median())
# Print the median of the dry_bulb_faren column for the time range
'2011-Apr': '2011-Jun'
print(df_clean.loc['2011-Apr': '2011-Jun', 'dry_bulb_faren'].median())
# Print the median of the dry_bulb_faren column for the month of
January
print(df_clean.loc['2011-Jan', 'dry_bulb_faren'].median())

72.0
78.0
48.0
```

```
# Downsample df_clean by day and aggregate by mean: daily_mean_2011
daily_mean_2011 = df_clean.resample('D').mean()
```

```
# Extract the dry_bulb_faren column from daily_mean_2011 using
.values: daily_temp_2011
daily_temp_2011 = daily_mean_2011['dry_bulb_faren'].values
```

```
# Downsample df_climate by day and aggregate by mean: daily_climate
daily_climate = df_climate.resample('D').mean()
```

```
# Extract the Temperature column from daily_climate using .reset_index():
daily_temp_climate = daily_climate.reset_index()['Temperature']
```

```
# Compute the difference between the two arrays and print the mean difference
difference = daily_temp_2011 - daily_temp_climate
print(difference.mean())
```

1.3301831870056477

```
# Using df_clean, when is sky_condition 'CLR'?
is_sky_clear = df_clean['sky_condition']=='CLR'
# Filter df_clean using is_sky_clear
sunny = df_clean.loc[is_sky_clear]
# Resample sunny by day then calculate the max
sunny_daily_max = sunny.resample('D').max()
# See the result
sunny_daily_max.head()
```

```
# Using df_clean, when does sky_condition contain 'OVC'?
is_sky_overcast = df_clean['sky_condition'].str.contains('OVC')
# Filter df_clean using is_sky_overcast
overcast = df_clean.loc[is_sky_overcast]
# Resample overcast by day then calculate the max
overcast_daily_max = overcast.resample("D").max()
# See the result
overcast_daily_max.head()
```

```
# From previous steps
is_sky_clear = df_clean['sky_condition']=='CLR'
sunny = df_clean.loc[is_sky_clear]
sunny_daily_max = sunny.resample('D').max()
is_sky_overcast = df_clean['sky_condition'].str.contains('OVC')
overcast = df_clean.loc[is_sky_overcast]
overcast_daily_max = overcast.resample('D').max()
# Calculate the mean of sunny_daily_max
sunny_daily_max_mean = sunny_daily_max.mean()
# Calculate the mean of overcast_daily_max
overcast_daily_max_mean = overcast_daily_max.mean()
# Print the difference (sunny minus overcast)
print(sunny_daily_max_mean - overcast_daily_max_mean)
```